

Section 1. Statistical Test

1.1 Which statistical test did you use to analyze the NYC subway data? Did you use a one-tail or a two-tail P value? What is the null hypothesis? What is your p-critical value?

Answer: statistical test: Mann-Whitney u-test (not a normal distribution) along with the mean of the two samples. **Two-tail test** given the fact that we don't know if the difference is greater or less than the average. **Null hypothesis:** There is no difference in Ridership on rainy days and ridership on days without rain. **P-critical** 0.025(two-tail),

1.2 Why is this statistical test applicable to the dataset? In particular, consider the assumptions that the test is making about the distribution of ridership in the two samples.

Answer: It's a non-parametric test; it doesn't assume the data is drawn from any particular probability distribution. The distributions of entries with and without rain are not normally distributed.

1.3 What results did you get from this statistical test? These should include the following numerical values: p-values, as well as the means for each of the two samples under test

Answer: P-value: 0,012499956, rain: 1105.4464, without rain: 1090.2788

1.4 What is the significance and interpretation of these results?

Answer: The two samples come from the same population, and are critical at 0,012499956 which falls in the critical value zone established ± 0.025 for the two tail test.

Section 2. Linear Regression

2.1 What approach did you use to compute the coefficients theta and produce prediction for `ENTRIESn_hourly` in your regression model:

1. **Gradient descent (as implemented in exercise 3.5)**
2. OLS using Statsmodels
3. Or something different?

2.2 What features (input variables) did you use in your model? Did you use any dummy variables as part of your features?

Answer: Rain, precipi, maxtempi, Hour, UNIT. Yes rain is a dummy variable 0 represent the absence of rain and 1 the presence of rain and UNIT takes 0 or 1 if a given station holds a record(unit_R001 = 1 if R001 shows up on that particular record for the UNIT column).

2.3 Why did you select these features in your model? We are looking for specific reasons that lead you to believe that the selected features will contribute to the predictive power of your model.

- Your reasons might be based on intuition. For example, response for fog might be: "I decided to use fog because I thought that when it is very foggy outside people might decide to use the subway more often."
- Your reasons might also be based on data exploration and experimentation, for example: "I used feature X because as soon as I included it in my model, it drastically improved my R^2 value."

Answer: I used rain because it's the most probable weather condition that would influence subway ridership and most common also(it can rain any time of the year.), maxtempi can indicate the temperature and what relationship temperature has on ridership. The hour variable helps pinpoint the exact moment in which ridership changes and if the same hour with different weather conditions varies on ridership. Precipitation would be an interesting factor that might answer questions like, does the amount of rain influence ridership or is it just the fact that it rains without any variability in the amount being of importance. UNIT has a strong significance on the coefficient of determination output and makes sense to keep track of ridership per train station.

2.4 What are the coefficients (or weights) of the non-dummy features in your linear regression model?

Answer:

Hour: 412.067

meantempi: -52.441

precipi: 21.883

2.5 What is your model's R^2 (coefficients of determination) value?

Answer: 0.463968977315, meaning that with this model we have explained 46% of the original data, and have a residual variability of 54%. But the R^2 value in this case is almost irrelevant, given the fact that what's being measured is the relationship between the predictors and response variable.

2.6 What does this R^2 value mean for the goodness of fit for your regression model? Do you think this linear model to predict ridership is appropriate for this dataset, given this R^2 value?

Answer: I think this coefficient of determination shows a value closer to 0, making it not appropriate, alpha value and features should be changed in order to get a better R^2 value.

Section 3. Visualization

Please include two visualizations that show the relationships between two or more variables in the NYC subway data.

Remember to add appropriate titles and axes labels to your plots. Also, please add a short description below each figure commenting on the key insights depicted in the figure.

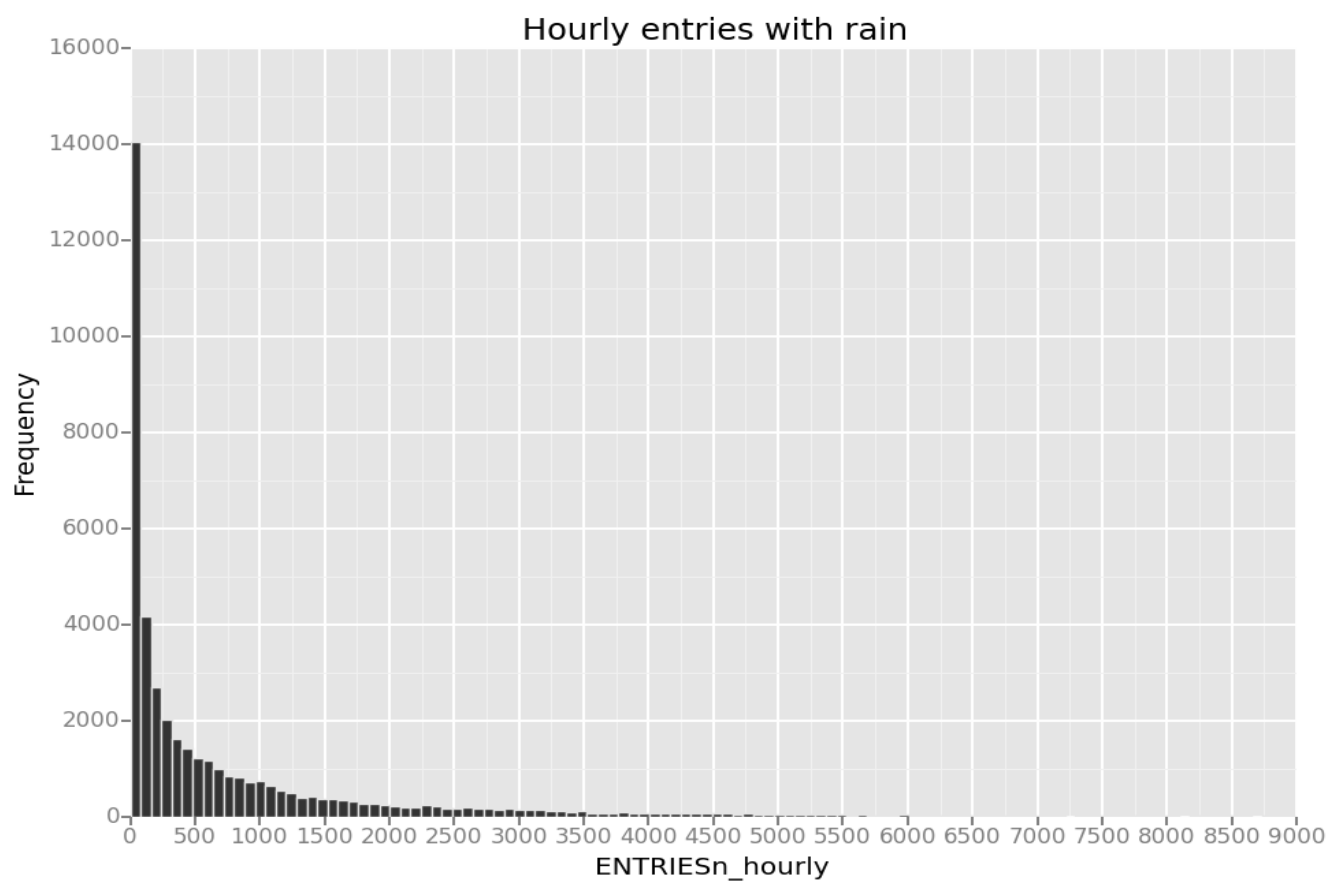
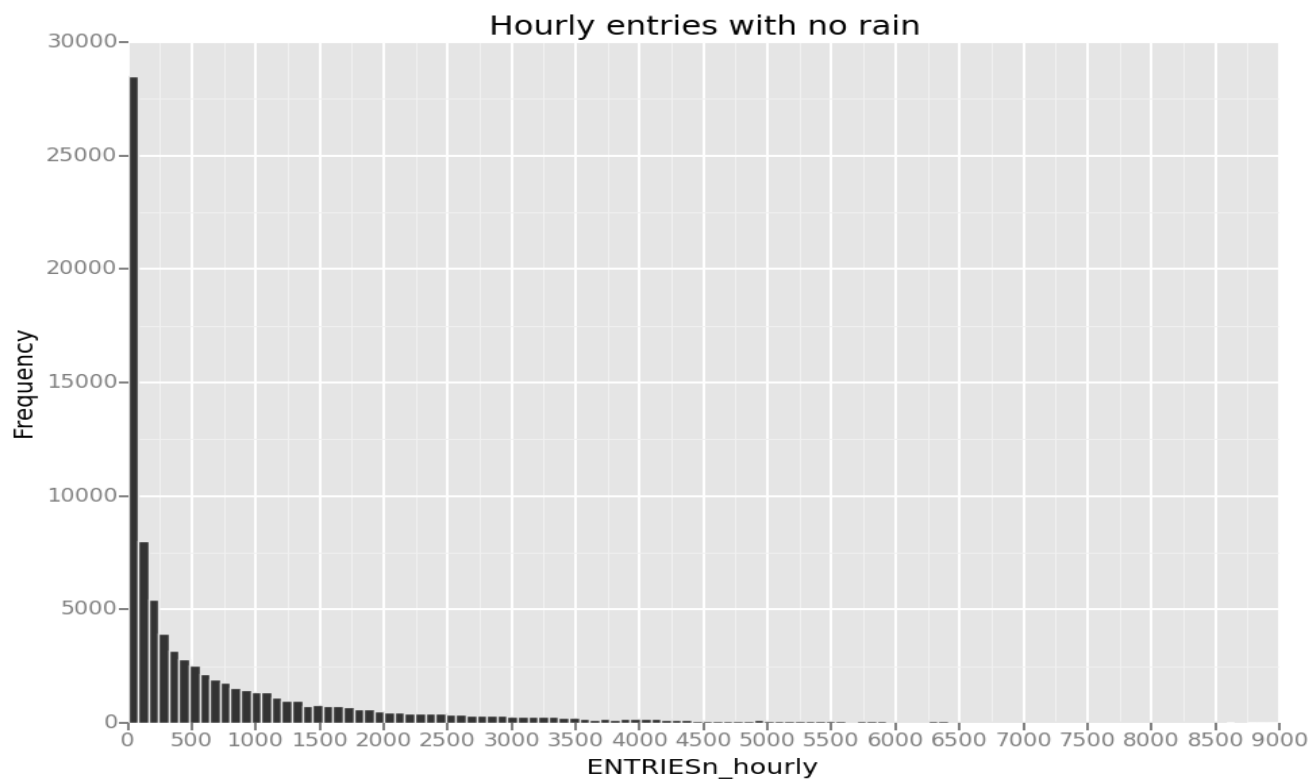
3.1 One visualization should contain two histograms: one of `ENTRIESn_hourly` for rainy days and one of `ENTRIESn_hourly` for non-rainy days.

You can combine the two histograms in a single plot or you can use two separate plots.

If you decide to use two separate plots for the two histograms, please ensure that the x-axis limits for both of the plots are identical. It is much easier to compare the two in that case.


For the histograms, you should have intervals representing the volume of ridership (value of `ENTRIESn_hourly`) on the x-axis and the frequency of occurrence on the y-axis. For example, each interval (along the x-axis), the height of the bar for this interval will represent the number of records (rows in our data) that have `ENTRIESn_hourly` that falls in this interval.

Remember to increase the number of bins in the histogram (by having larger number of bars). The default bin width is not sufficient to capture the variability in the two samples.

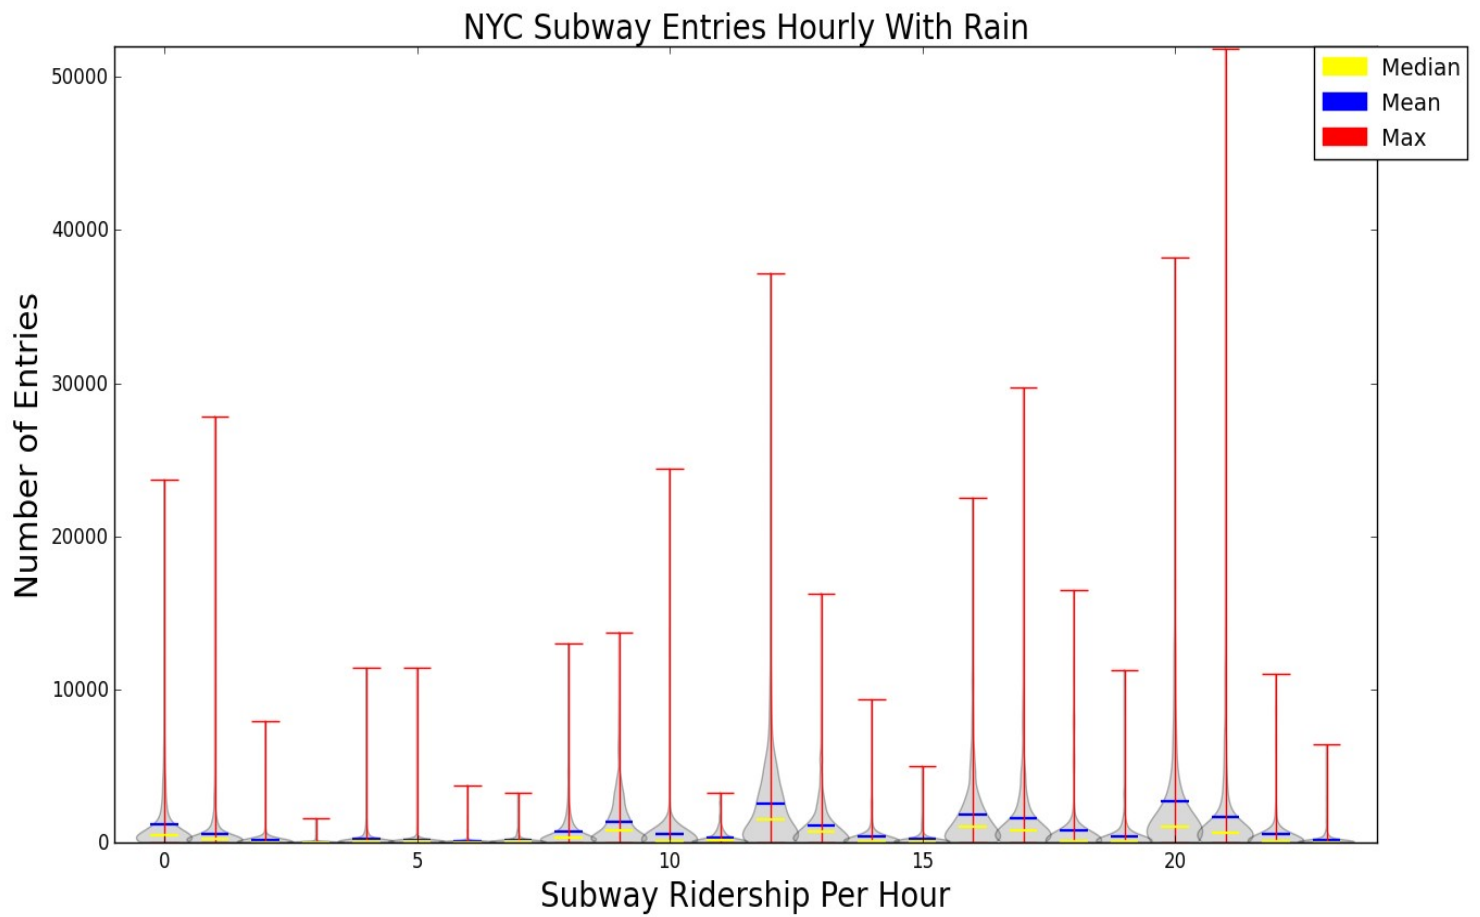


3.2 One visualization can be more freeform. You should feel free to implement something that we discussed in class (e.g., scatter plots, line plots) or attempt to implement something more advanced if you'd like.

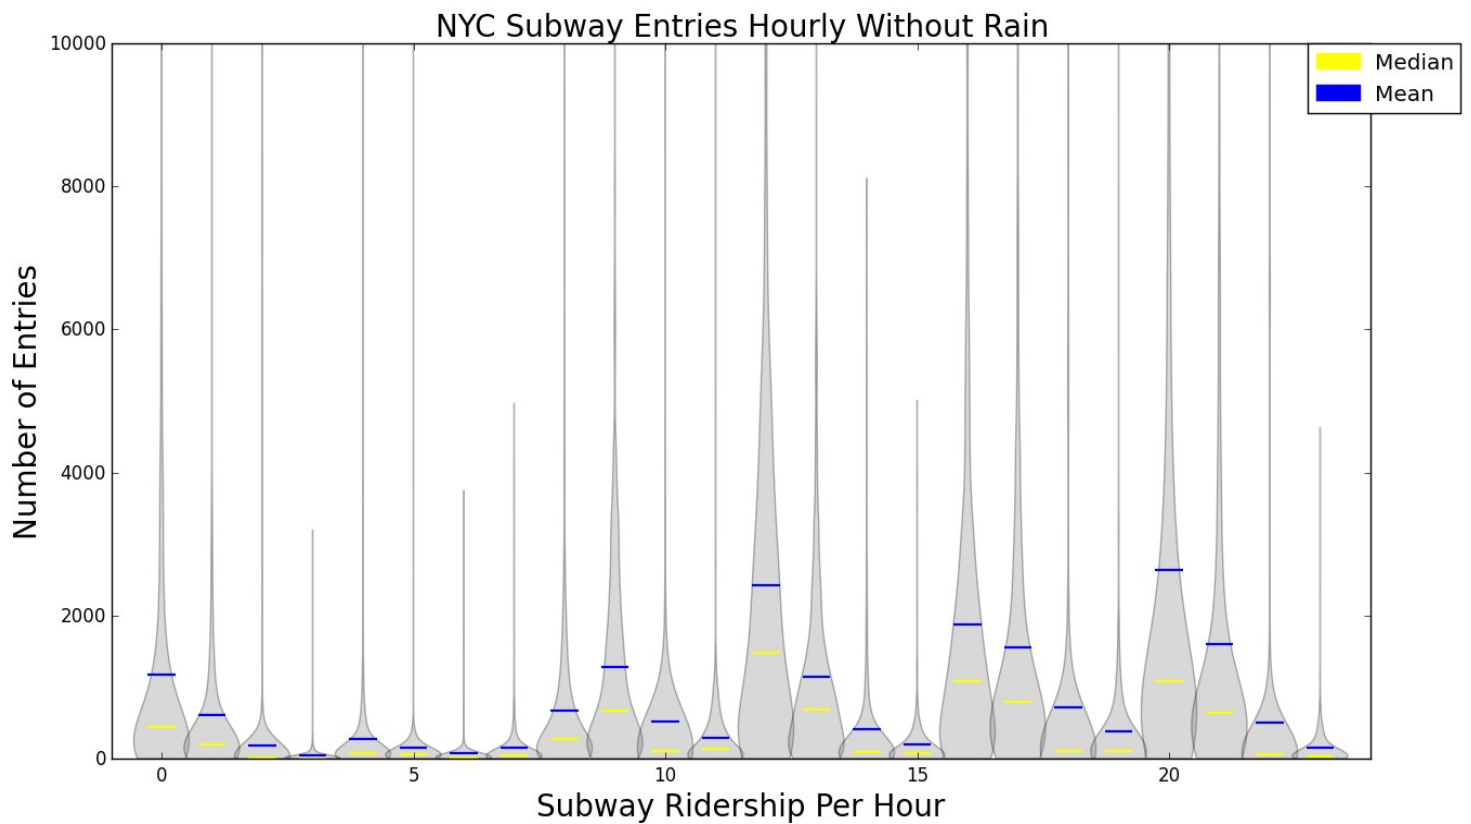
Description: In this graphs I intend to show the difference between entries in the NYC Subway in rainy and normal days. The goal it's to detect variability in each hour of the day.




This plot shows ridership on a day without rain, showing the max values for each hour.



This plot shows ridership on a day with rain, showing the max values for each hour.



This plot shows ridership on a day without rain, showing the mean and median with more detail.



This plot shows ridership on a day with rain, showing the mean and median with more detail.

Section 4. Conclusion

Please address the following questions in detail. Your answers should be 1-2 paragraphs long.

4.1 From your analysis and interpretation of the data, do more people ride the NYC subway when it is raining or when it is not raining?

Answer: From my analysis I can gather a couple of conclusions, first the null hypothesis was rejected, because there is a significant difference between the mean ridership on rainy vs non rainy days. Despite the fact that the average on rainy days is higher than the one on days without rain it's hard to make a prediction on whether ridership is truly higher given that the number of records in each category is very unbalanced and the higher number of 0 entries on the days without rain brings down the average, but my gathering and interpretation leads me into believing that ridership for the NYC subway is higher on rainy days (Rainy days mean: 1105.4464 – No rain days mean: 1090.2788).

4.2 What analyses lead you to this conclusion? You should use results from both your statistical tests and your linear regression to support your analysis.

Answer: Mann-Whitney u-test (not a normal distribution) along with the mean of the two samples. P-value: 0.012499956, rain: 1105.4464, without rain: 1090.2788. This shows higher ridership average on rainy days compared to an example of days without rain from the same distribution. R^2 (coefficients of determination) value 0.463968977315 demonstrates a working model that has room for improvement.

Section 5. Reflection

Please address the following questions in detail. Your answers should be 1-2 paragraphs long.

5.1 Please discuss potential shortcomings of the methods of your analysis, including:

1. Dataset,

Answer: The dataset has a couple of potential shortcomings. First, there appears to be many missing values from the Entries_hourly column that have been filled in with random zeros. Second, there are many different types of value representations, dates, float, ints and even strings which could make it hard to work on some occasions.

2. Analysis, such as the linear regression model or statistical test.

Answer: The Mann-Whitney Test it's a non-parametric test that lets us bypass knowledge on the distribution of the data, but comes at a price of lower power, meaning, if there is a significant difference between the two samples it will be harder to find than in a parametric test. The linear regression model will only allow insights into linear relationships and right now might be wrongly influenced by outliers.

Bibliography

Linear regression:

source: University of duke

url: https://stat.duke.edu/courses/Fall00/sta103/lecture_notes/multregr.pdf

Mann-Whitney U Test:

source: Laerd Statitics

url: <https://statistics.laerd.com/spss-tutorials/mann-whitney-u-test-using-spss-statistics.php>

Ggplot

source Yhat

url:<https://ggplot.yhathq.com/>
